9. Bytecode and Virtual Machines

Oscar Nierstrasz

Original material prepared by Adrian Lienhard and Marcus Denker

Roadmap

> Introduction
> Bytecode
> The heap
> Interpreter
> Automatic memory management
> Threading System
> Optimizations

Birds-eye view

A virtual machine is an abstract computing architecture supporting a programming language in a hardware-independent fashion
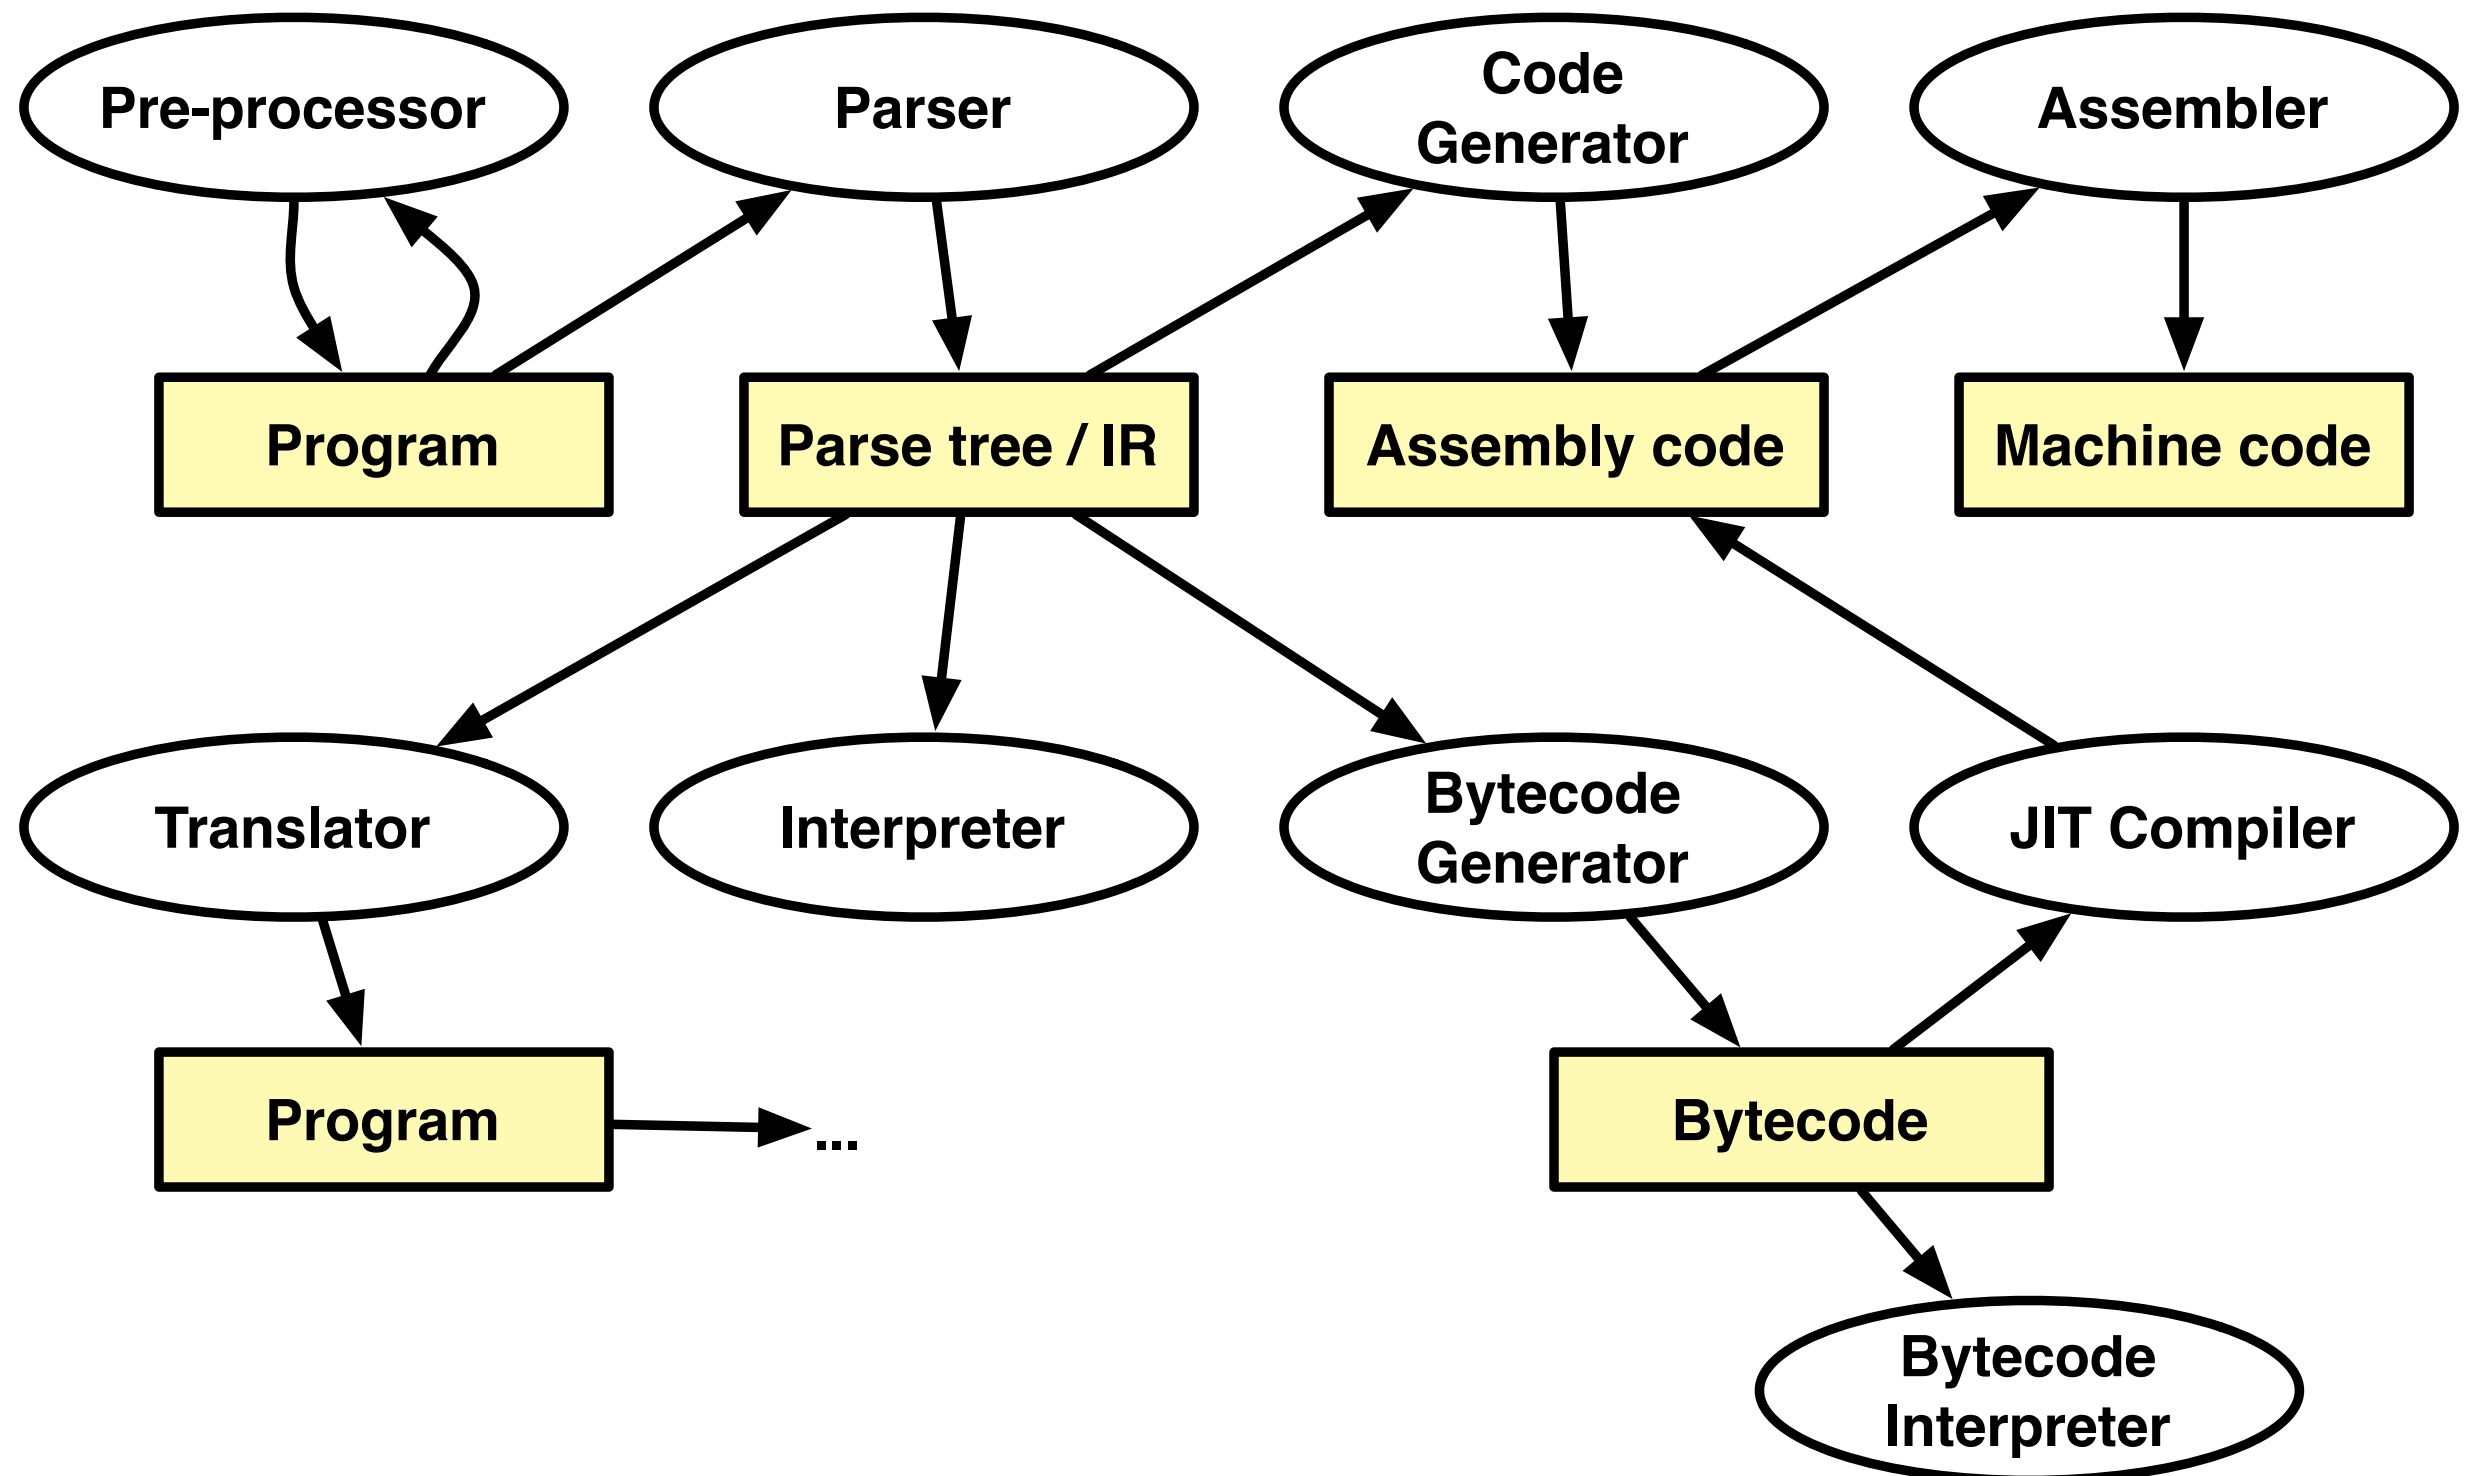
Z1, 1938

Roadmap

Implementing a Programming Language

How are VMs implemented?

Typically using an *efficient and portable language* such as C, C++, or assembly code

Pharo VM platform-independent part written in *Slang:*
- subset of Smalltalk, translated to C
- core: 600 methods or 8k LOC in Slang
- Slang allows one to simulate VM in Smalltalk

Main Components of a VM

The heap

The interpreter

Automatic memory management

The threading System

Pros and Cons of the VM Approach

Pros

> Platform independence of application code
  "Write once, run anywhere"

> Simpler programming model

> Security

> Optimizations for different hardware architectures

Cons

> Execution overhead

> Not suitable for system programming

Roadmap

> Introduction
> **Bytecode**
> The heap
> Interpreter
> Automatic memory management
> Threading System
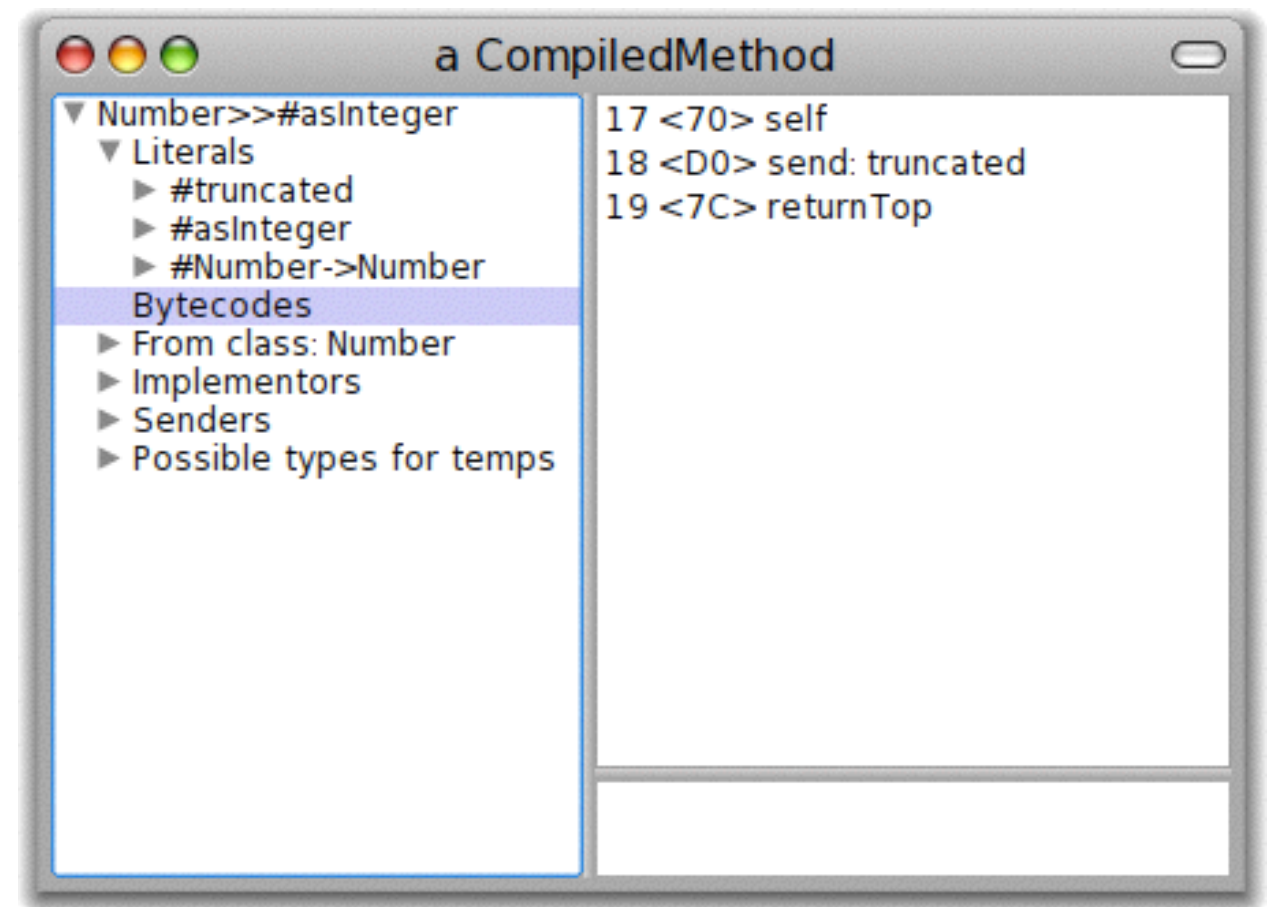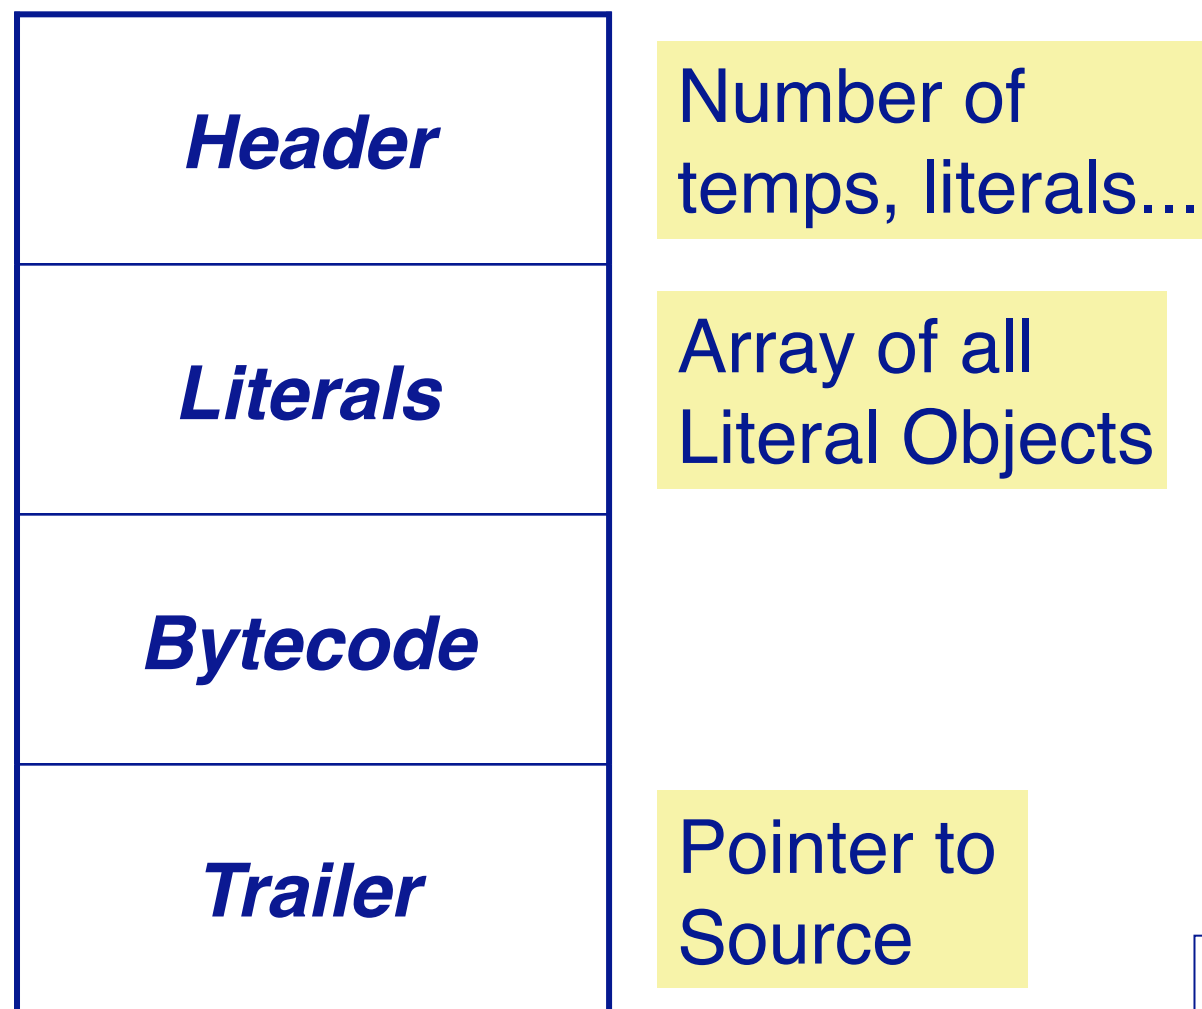> Optimizations

Reasons for working with Bytecode

> Generating Bytecode
— Implementing compilers for other languages
— Experimentation with new language features

> Parsing and Interpretation:
— Analysis (e.g., `self` and `super` sends)
— Decompilation (for systems without source)
— Printing of bytecode
— Interpretation: `Debugger`, `Profiler`

The Pharo Virtual Machine

> Virtual machine provides a virtual processor
—Bytecode: The "machine-code" of the virtual machine

> Smalltalk (like Java): Stack machine
—easy to implement interpreters for different processors
—most hardware processors are register machines

> Pharo VM: Implemented in *Slang*
—Slang: Subset of Smalltalk. ("C with Smalltalk Syntax")
—Translated to C

Bytecode in the CompiledMethod

> CompiledMethod format:

Header	Number of temps, literals...
Literals	Array of all Literal Objects
Bytecode	
Trailer	Pointer to Source

Integer>>#timesRepeat:.

```
a CompiledMethod

▼ Number>>#asInteger          17 <70> self
  ▼ Literals                  18 <D0> send: truncated
    ► #truncated              19 <7C> returnTop
    ► #asInteger
    ► #Number->Number
    Bytecodes
  ► From class: Number
  ► Implementors
  ► Senders
  ► Possible types for temps
```

```
(Number>>#asInteger) inspect
```

```
(Number methodDict at: #asInteger) inspect
```

Bytecodes: Single or multibyte

> Different forms of bytecodes:

—Single bytecodes:

– *Example: 112: push self*

—Groups of similar bytecodes

Type	*Offset*
4 bits	4 bits

– *16: push temp 1*

– *17: push temp 2*

– *up to 31*

—Multibyte bytecodes

– *Problem: 4 bit offset may be too small*

– *Solution: Use the following byte as offset*

– *Example: Jumps need to encode large jump offsets*

Example: Number>>asInteger

> Smalltalk code:

```
Number>>asInteger
    "Answer an Integer nearest
    the receiver toward zero."

    ^self truncated
```

> Symbolic Bytecode

```
17 <70> self
18 <D0> send: truncated
19 <7C> returnTop
```

Example: Step by Step

> `17 <70> self`

  — The receiver (self) is pushed on the stack

> `18 <D0> send: truncated`

  — Bytecode 208: send literal selector 1

  — Get the selector from the first literal

  — start message lookup in the class of the object that is on top of the stack

  — result is pushed on the stack

> `19 <7C> returnTop`

  — return the object on top of the stack to the calling method

Pharo Bytecode

> 256 Bytecodes, four groups:

— Stack Bytecodes
- *Stack manipulation: push / pop / dup*

— Send Bytecodes
- *Invoke Methods*

— Return Bytecodes
- *Return to caller*

— Jump Bytecodes
- *Control flow inside a method*

Stack Bytecodes

> Push values on the stack
  —e.g., temps, instVars, literals
  —e.g: 16 - 31: push instance variable

> Push Constants
  —False/True/Nil/1/0/2/-1

> Push `self`, `thisContext`

> Duplicate top of stack

> Pop

Sends and Returns

> Sends: receiver is on top of stack
  —Normal send
  —Super Sends
  —Hard-coded sends for efficiency, e.g. +, –

> Returns
  —Return top of stack to the sender
  —Return from a block
  —Special bytecodes for return `self`, `nil`, `true`, `false` (for efficiency)

Jump Bytecodes

> Control Flow inside one method
  —Used to implement control-flow efficiently
  —Example:

```
^ 1<2 ifTrue: ['true']
```

```
17 <76> pushConstant: 1
18 <77> pushConstant: 2
19 <B2> send: <
20 <99> jumpFalse: 23
21 <20> pushConstant: 'true'
22 <90> jumpTo: 24
23 <73> pushConstant: nil
24 <7C> returnTop
```
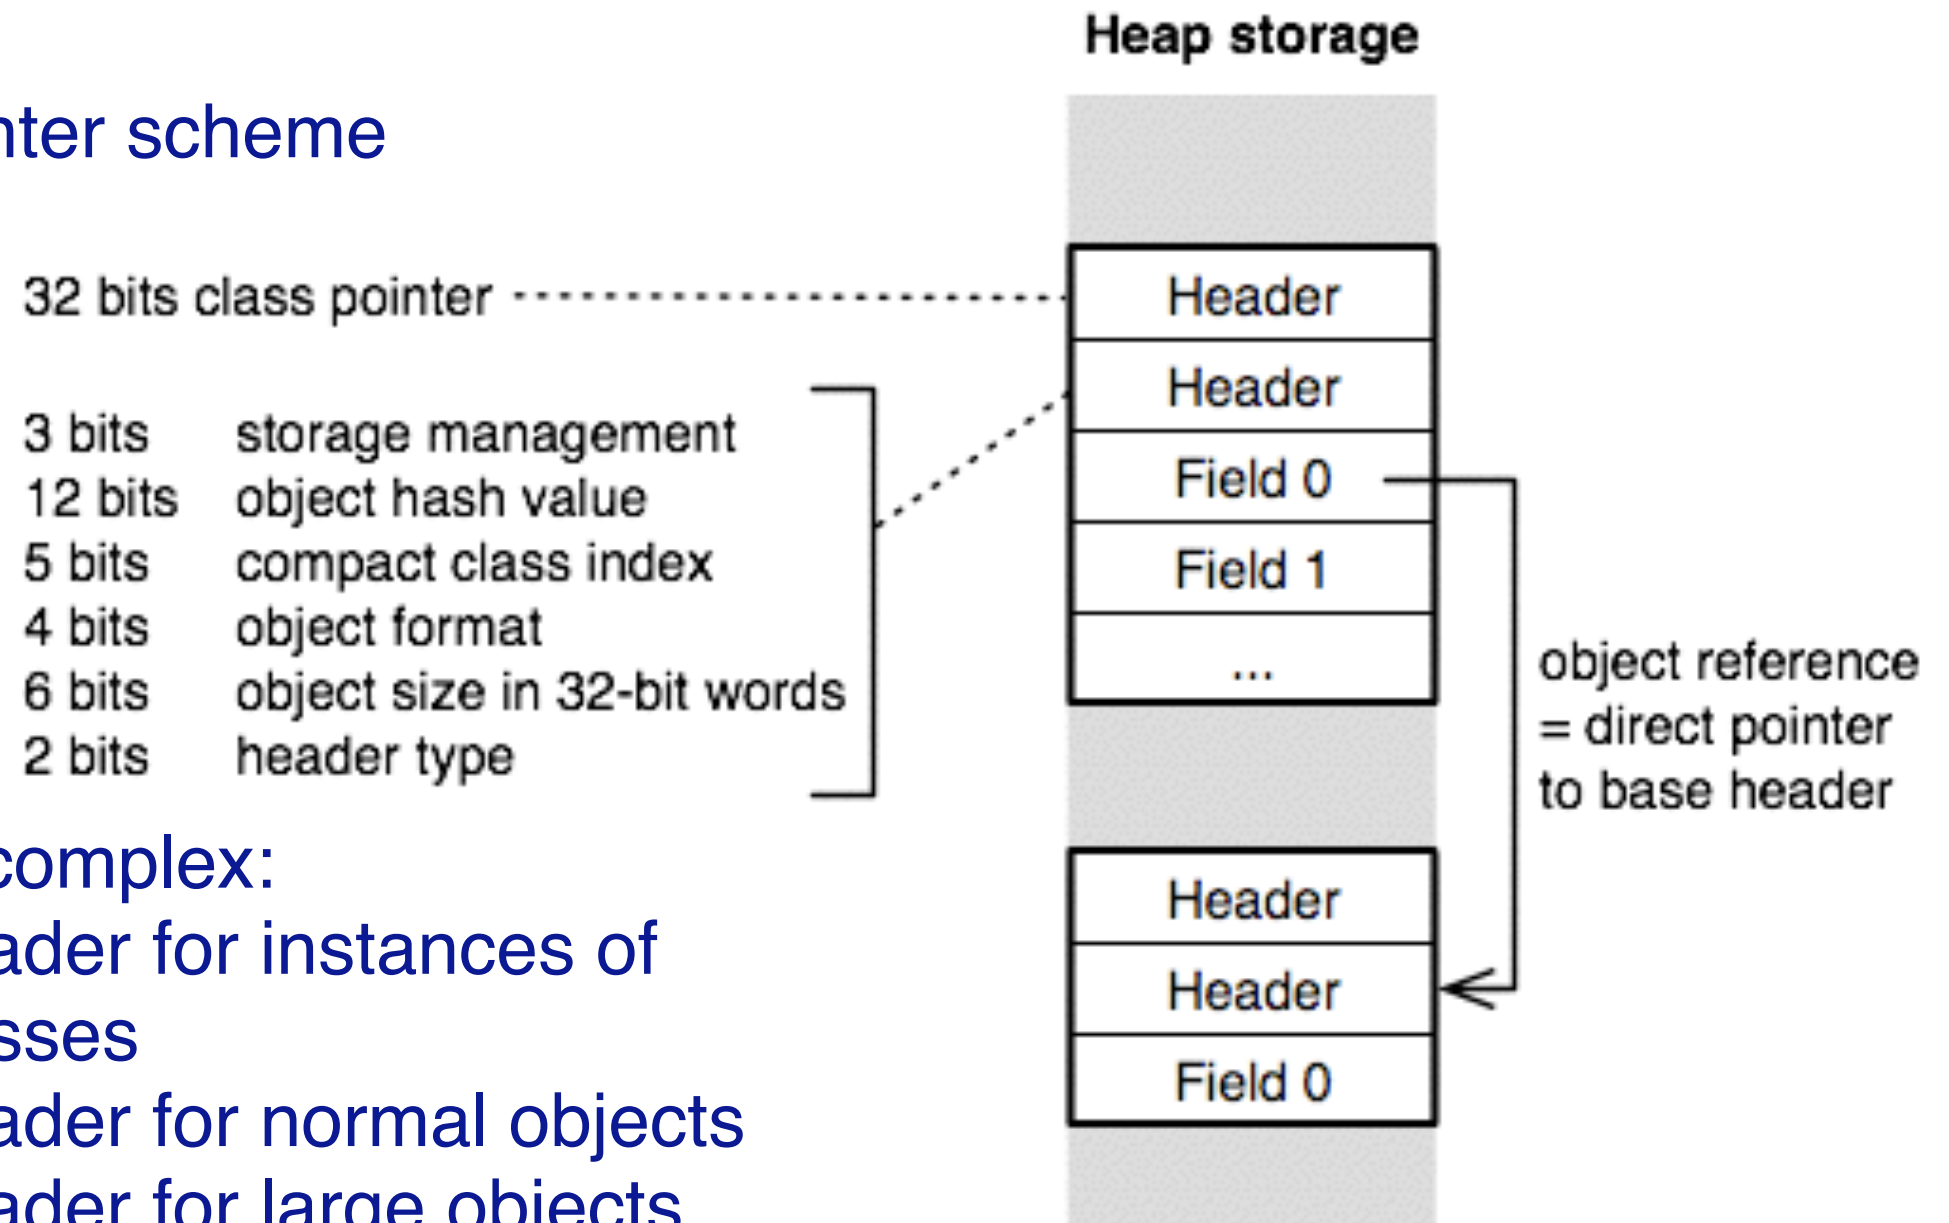
Roadmap

> Introduction
> Bytecode
> **The heap**
> Interpreter
> Automatic memory management
> Threading System
> Optimizations

Object Memory Layout

32-bit direct-pointer scheme

Heap storage

32 bits class pointer ···→ Header

3 bits	storage management
12 bits	object hash value
5 bits	compact class index
4 bits	object format
6 bits	object size in 32-bit words
2 bits	header type

Header

Header

Field 0

Field 1

...

object reference
= direct pointer
to base header

Header

Header

Field 0

Reality is more complex:
- 1-word header for instances of compact classes
- 2-word header for normal objects
- 3-word header for large objects

Different Object Formats

> fixed pointer fields

> indexable types:
  – indexable pointer fields (e.g., Array)
  – indexable weak pointer fields (e.g., WeakArray)
  – indexable word fields (e.g., Bitmap)
  – indexable byte fields (e.g., ByteString)

Object format (4bit)
     0   no fields
     1   fixed fields only
     2   indexable pointer fields only
     3   both fixed and indexable pointer fields
     4   both fixed and indexable weak fields
     6   indexable word fields only
     8-11   indexable byte fields only
     12-15 ...

Iterating Over All Objects in Memory

```
"Answer the first object on the heap"
anObject someObject

"Answer the next object on the heap"
anObject nextObject
```

Excludes small integers!

```
SystemNavigation>>allObjectsDo: aBlock
    | object endMarker |
    object := self someObject.
    endMarker := Object new.
    [endMarker == object]
        whileFalse: [aBlock value: object.
            object := object nextObject]
```

```
|count|
count := 0.
SystemNavigation default allObjectsDo:
    [:anObject | count := count + 1].
count
```

529468

Roadmap

> Introduction
> Bytecode
> The heap
> **Interpreter**
> Automatic memory management
> Threading System
> Optimizations

Stack vs. Register VMs

VM provides a virtual processor that interprets bytecode instructions

Stack machines
- Smalltalk, Java and most other VMs
- Simple to implement for different hardware architectures
- Very compact code

Register machines
- Potentially faster than stack machines
- Only few register VMs, e.g., Parrot VM (Perl6)

Interpreter State and Loop

Interpreter state
- instruction pointer (ip): points to current bytecode
- stack pointer (sp): topmost item in the operand stack
- current active method or block context
- current active receiver and method

Interpreter loop
1. branch to appropriate bytecode routine
2. fetch next bytecode
3. increment instruction pointer
4. execute the bytecode routine
5. return to 1.

Method Contexts
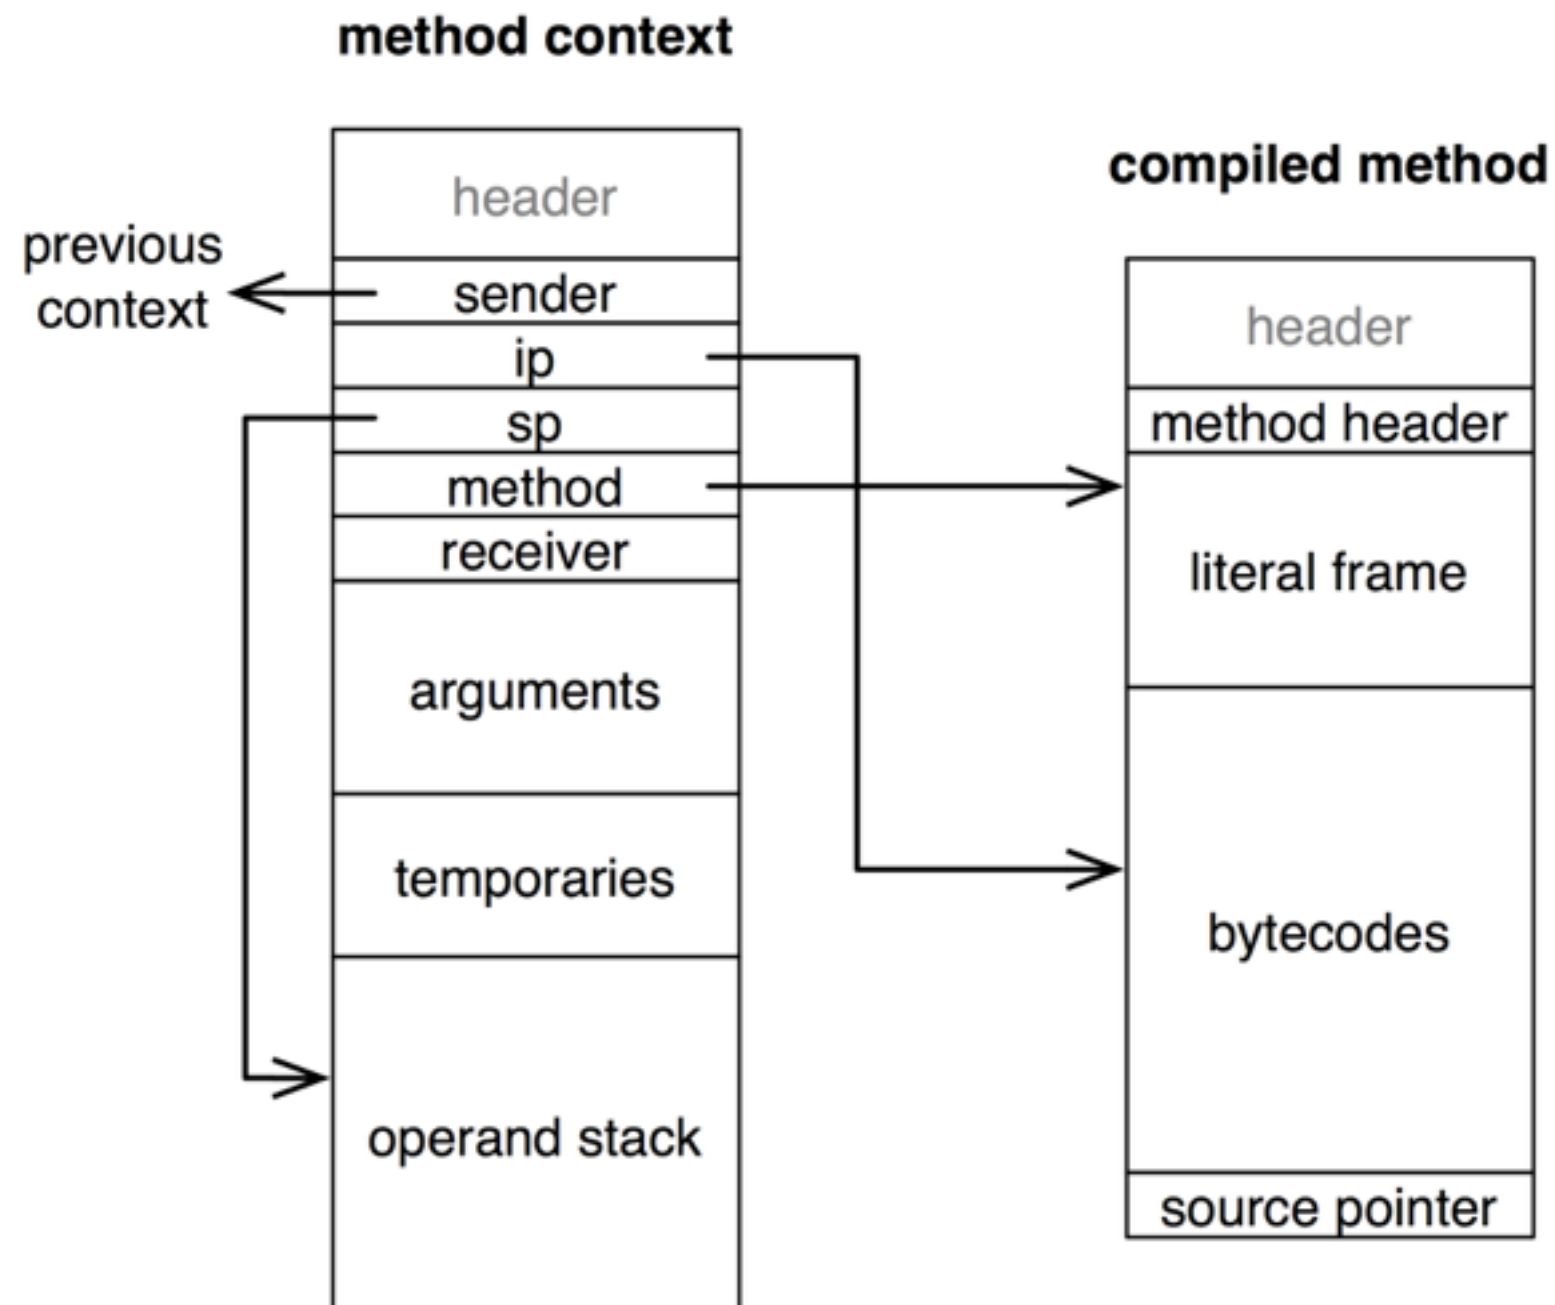
method header:
– primitive index
– number of args
– number of temps
– large context flag
– number of literals

Stack Manipulating Bytecode Routine

Example: bytecode <70> self

```
Interpreter>>pushReceiverBytecode
   self fetchNextBytecode.
   self push: receiver


Interpreter>>push: anObject
   sp := sp + BytesPerWord.
   self longAt: sp put: anObject
```

Stack Manipulating Bytecode Routine

Example: bytecode `<01>` `pushRcvr: 1`

```
Interpreter>>pushReceiverVariableBytecode
  self fetchNextBytecode.
  self pushReceiverVariable: (currentBytecode bitAnd: 16rF)

Interpreter>>pushReceiverVariable: fieldIndex
  self push: (
    self fetchPointer: fieldIndex ofObject: receiver)

Interpreter>>fetchPointer: fieldIndex ofObject: oop
  ^ self longAt: oop + BaseHeaderSize + (fieldIndex * BytesPerWord)
```

Message Sending Bytecode Routine

Example: bytecode `<E0> send: hello`

1. find selector, receiver and its class

2. lookup message in the method dictionary of the class

3. if method not found, repeat this lookup in successive superclasses; if superclass is nil, instead send #doesNotUnderstand:

4. create a new method context and set it up

5. activate the context and start executing the instructions in the new method

Message Sending Bytecode Routine

Example: bytecode `<E0>` `send: hello`

```
Interpreter>>sendLiteralSelectorBytecode
  selector := self literal: (currentBytcode bitAnd: 16rF).
  argumentCount := ((currentBytcode >> 4) bitAnd: 3) - 1.
  rcvr := self stackValue: argumentCount.
  class := self fetchClassOf: rcvr.
  self findNewMethod.
  self executeNewMethod.
  self fetchNewBytecode
```

This routine (bytecodes 208-255) can use any of the first 16 literals and pass up to 2 arguments

```
E0(hex) = 224(dec)
   = 1110 0000(bin)

E0 AND F = 0
   => literal frame at 0

((E0 >> 4) AND 3) - 1 = 1
   => 1 argument
```

Primitives

Primitive methods trigger a VM routine
and are executed without a new
method context unless they fail

```
ProtoObject>>nextObject
    <primitive: 139>
    self primitiveFailed
```

> Improve performance (arithmetics, at:, at:put:, ...)
> Do work that can only be done in VM (new object creation, process manipulation, become, ...)
> Interface with outside world (keyboard input, networking, ...)
> Interact with VM plugins (named primitives)

Roadmap

> Introduction
> Bytecode
> The heap
> Interpreter
> **Automatic memory management**
> Threading System
> Optimizations

Automatic Memory Management

Tell when an object is no longer used and then recycle the memory

Challenges
– Fast allocation
– Fast program execution

– Small predictable pauses
– Scalable to large heaps
– Minimal space usage

Main Approaches

> 1. Reference Counting
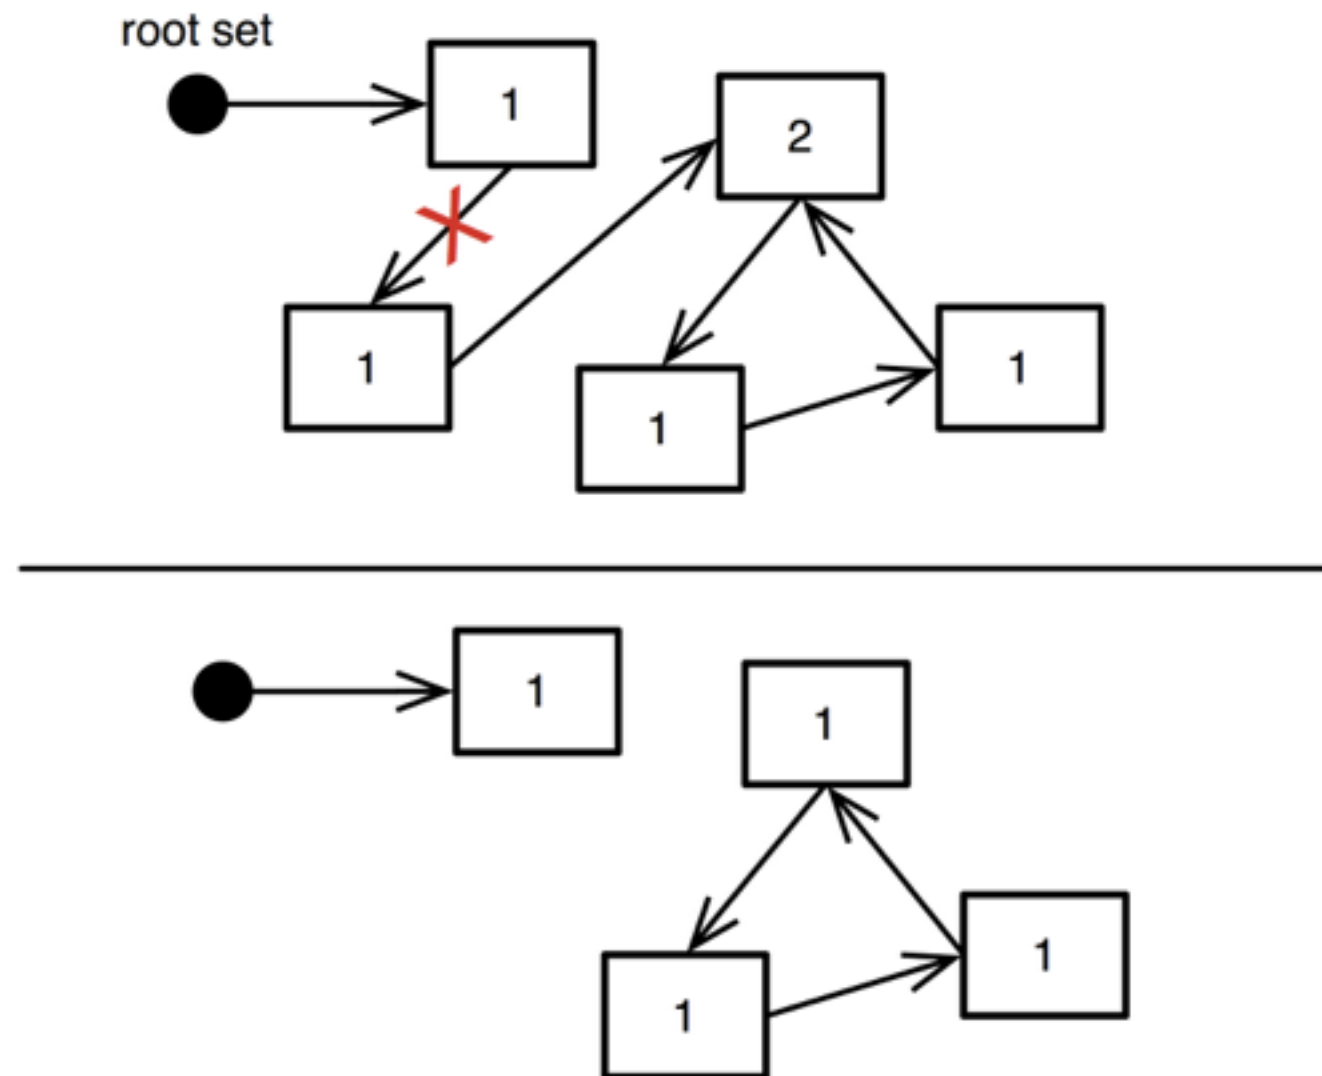
> 2. Mark and Sweep

Reference Counting GC

Idea

> For each store operation increment count field in header of newly stored object

> Decrement if object is overwritten

> If count is 0, collect object and decrement the counter of each object it pointed to

Problems

> Run-time overhead of counting (particularly on stack)

> Inability to detect cycles (need additional GC technique)

# Reference Counting GC

Mark and Sweep GC

Idea

> Suspend current process

> Mark phase: trace each accessible object leaving a mark in the object header (start at known root objects)

> Sweep phase: all objects with no mark are collected

> Remove all marks and resume current process

Problems

> Need to "stop the world"

> Slow for large heaps ➜ generational collectors

> Fragmentation ➜ compacting collectors

Mark and Sweep GC
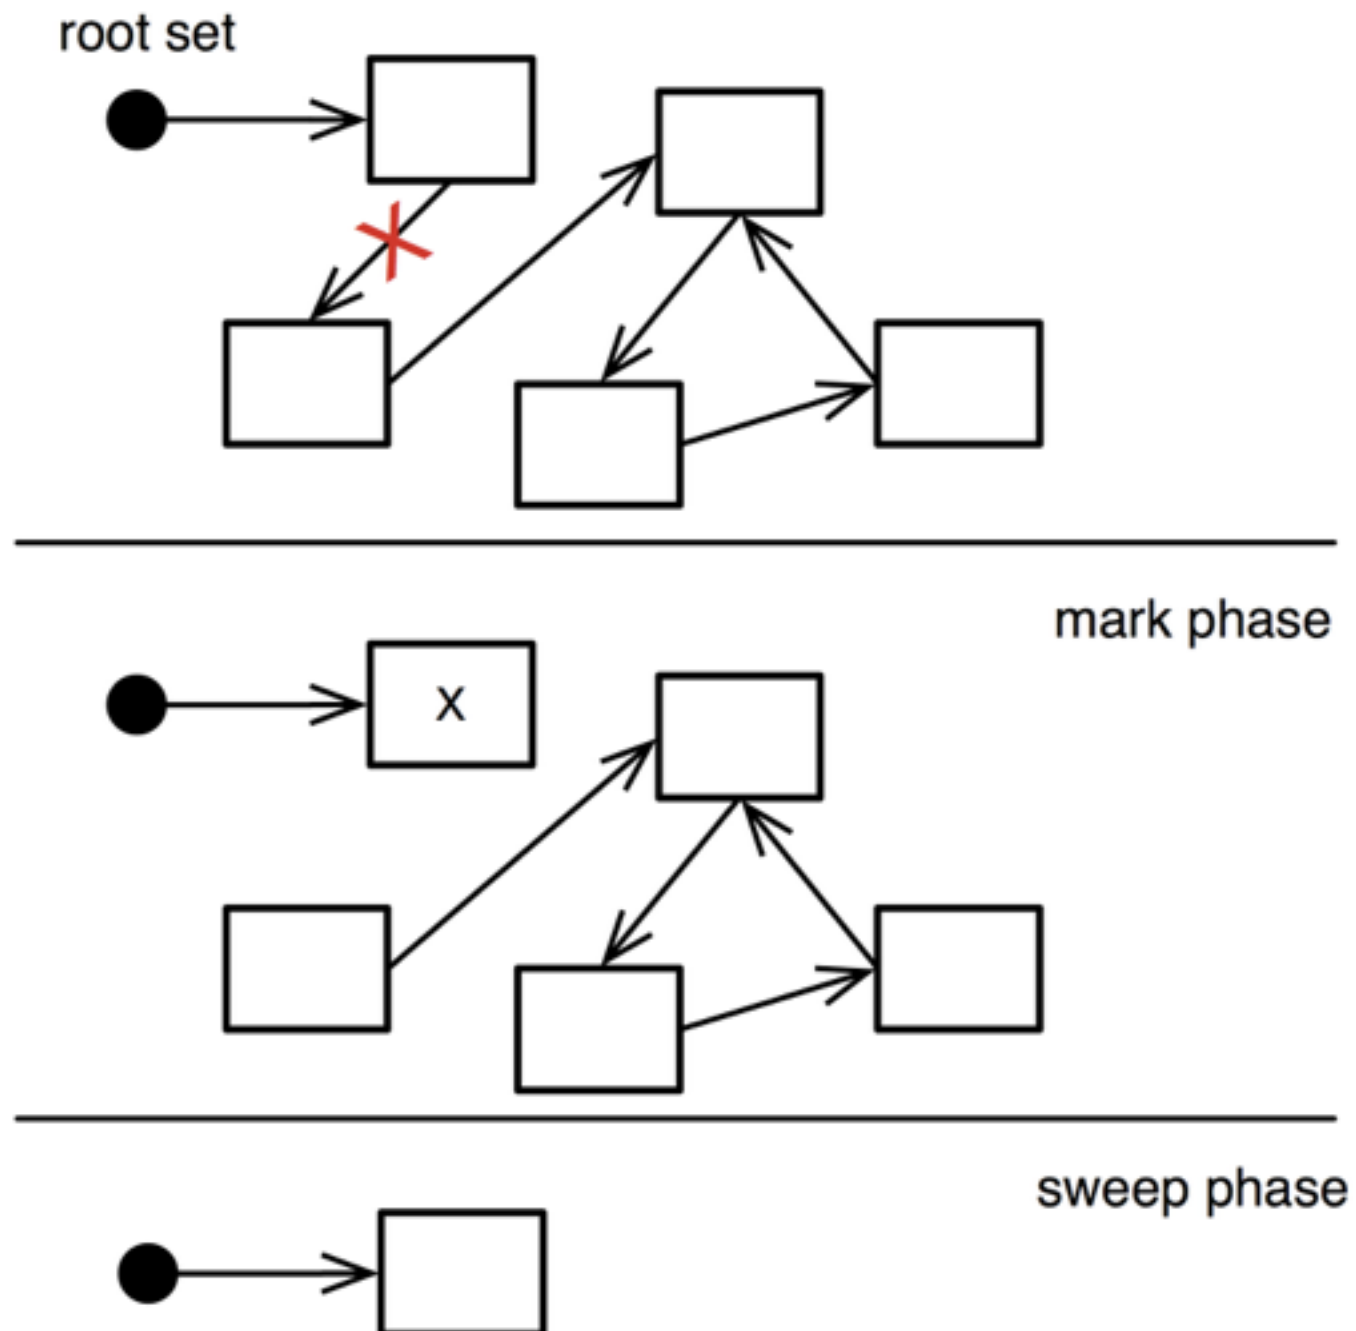
Generational Collectors

Most new objects live very short lives; most older objects live forever [Ungar 87]

Idea

> Partition objects into generations
> Create objects in young generation
> Tenuring: move live objects from young to old generation
> Incremental GC: frequently collect young generation (very fast)
> Full GC: infrequently collect young+old generation (slow)

Difficulty

> Need to track pointers from old to new space

Generational Collectors: Remembered Set

Write barrier: remember objects with <mark>old-young pointers:</mark>

> On each store check whether stored object (object2) is young and storer (object1) is old

```
object1.f := object2
```

> If true, add storer to remembered set

> When marking young generation, use objects in remembered set as additional roots
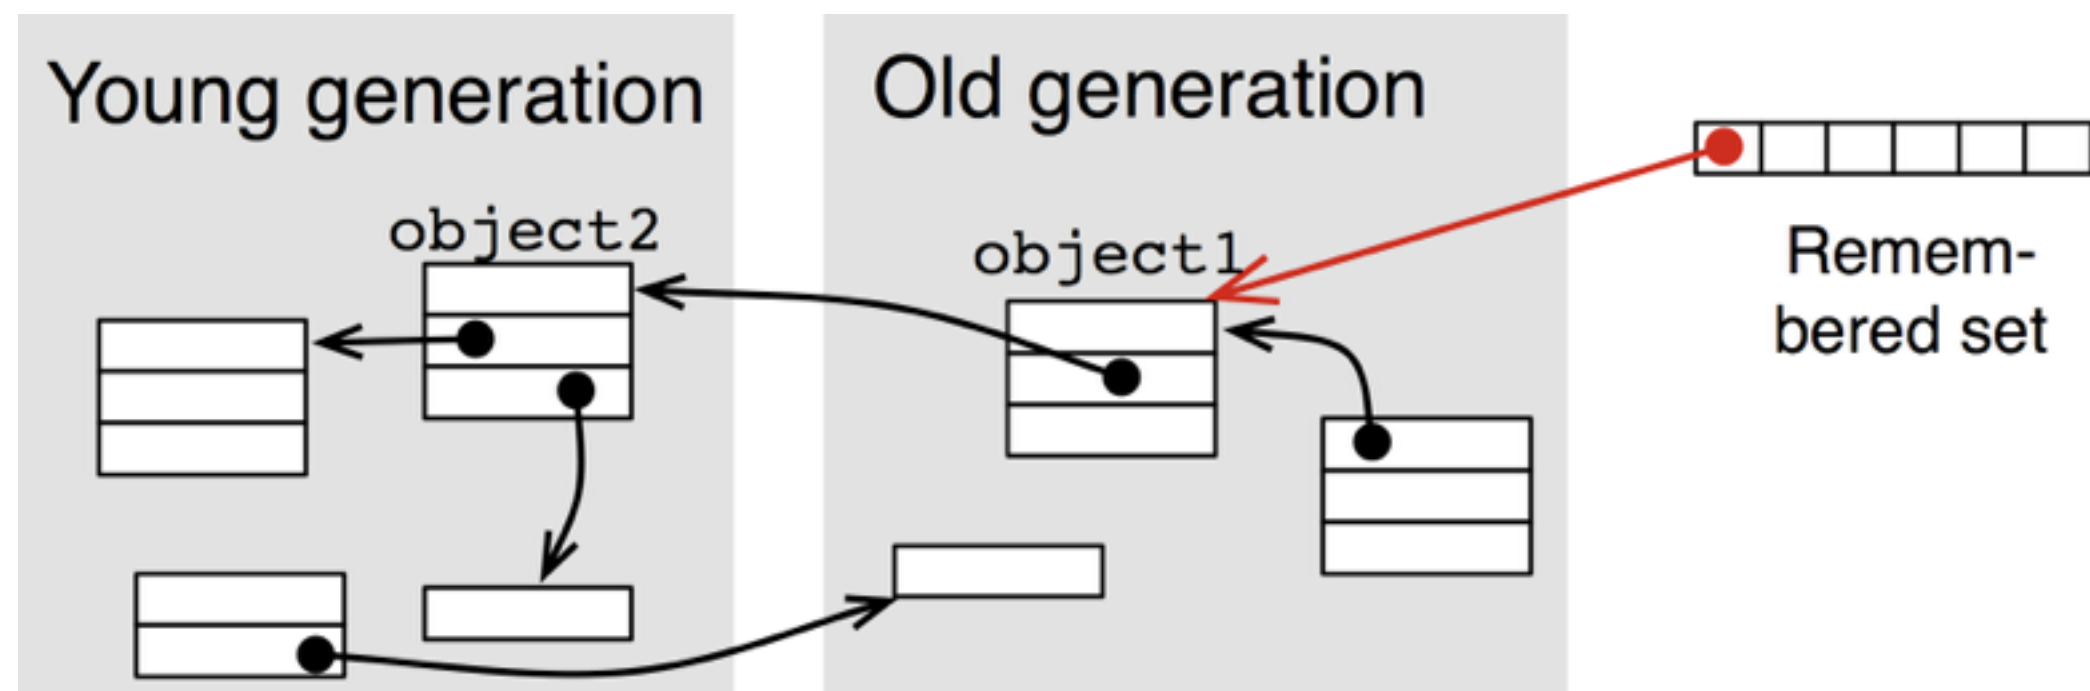
Compacting Collectors

Idea

> During the sweep phase all live objects are ==packed to the beginning of the heap==

> Simplifies allocation since free space is in one contiguous block

Challenge

> Adjust all pointers of moved objects

– object references on the heap

– pointer variables of the interpreter!

The Pharo GC

Pharo: <u>mark and sweep</u> <u>compacting collector</u> with <u>two generations</u>

> Cooperative, i.e., not concurrent
> Single threaded

When Does the GC Run?

– Incremental GC on allocation count or memory needs
– Full GC on memory needs
– Tenure objects if survivor threshold exceeded

```
"Tenure when more than this many objects survive"
Smalltalk vm tenuringThreshold                    2000
```

VM Memory Statistics

```
Smalltalk vm statisticsReport
```

```
'uptime                          0h8m6s
memory                           53,102,204 bytes
          old                         48,996,624 bytes (92.300000000000001%)
          young                  -2,105,420 bytes (-4.0%)
          used                   46,891,204 bytes (88.300000000000001%)
          free                   6,211,000 bytes (11.700000000000001%)
GCs                                   580 (839ms between GCs)
          full                        7 totalling 316ms (0.1% uptime), avg 45.1ms
          incr                   573 totalling 338ms (0.1% uptime), avg 0.600000000000001ms
          tenures                137 (avg 4 GCs/tenure)
Since last view    25 (654ms between GCs)
          uptime                 16.3s
          full                        0 totalling 0ms (0.0% uptime)
          incr                   25 totalling 13ms (0.1% uptime), avg 0.5ms
          tenures                1 (avg 25 GCs/tenure)
'
```

Memory System API

```
"Force GC"
Smalltalk garbageCollectMost.
Smalltalk garbageCollect.

"Is object young?"
Smalltalk isYoung: anObject.

"Various settings and statistics"
Smalltalk vm getParameters.

"Do an incremental GC after this many allocations"
Smalltalk vm allocationsBetweenGC: 4000.

"Tenure when more than this many objects survive the GC"
Smalltalk vm tenuringThreshold: 2000.

"Grow/shrink headroom"
Smalltalk vm parameterAt: 25 put: 4*1024*1024.
Smalltalk vm parameterAt: 24 put: 8*1024*1024.
```
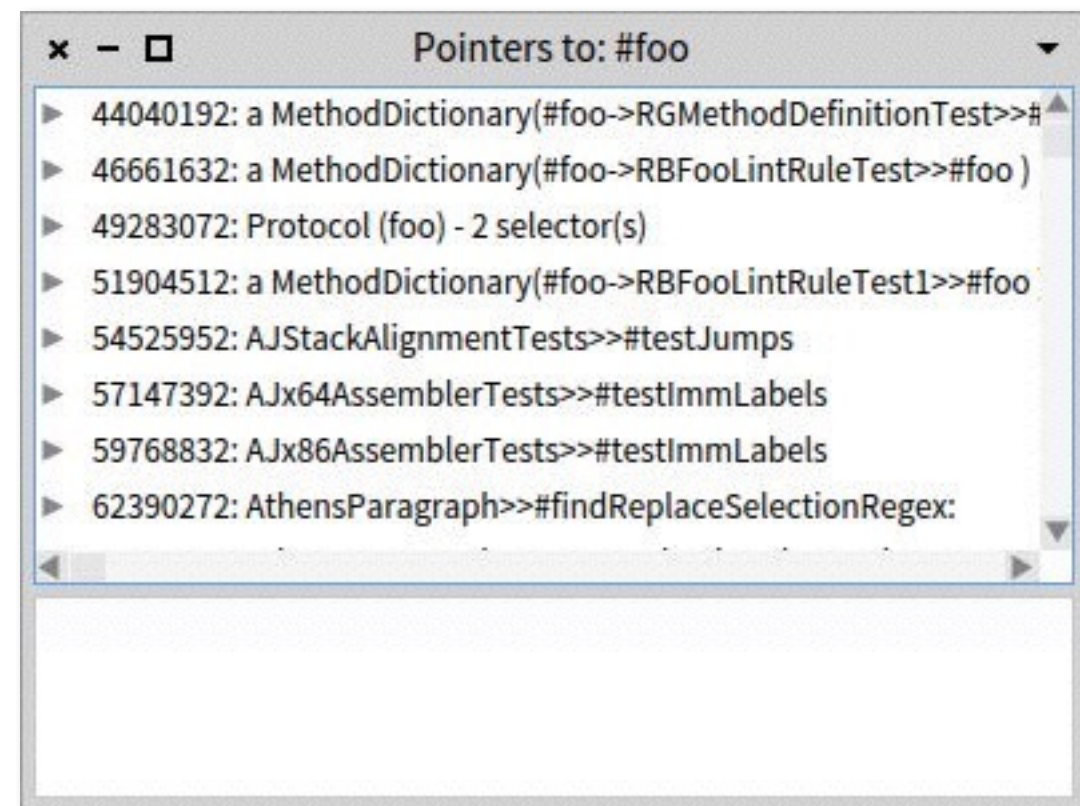
Finding Memory Leaks

I have objects that do not get collected. What's wrong?

– maybe object is just not GCed *yet* (force a full GC!)
– find the objects and then explore who references them

The pointer finder finds a path
from a root to some object

```
EyePointerExplorer openOn: #foo
```

Roadmap

> Introduction
> Bytecode
> The heap
> Interpreter
> Automatic memory management
> **Threading System**
> Optimizations

Threading System

Multithreading is the ability to create concurrently running "processes"

Non-native threads (*green threads*)
- Only one native thread used by the VM
- Simpler to implement and easier to port

Native threads
- Using the native thread system provided by the OS
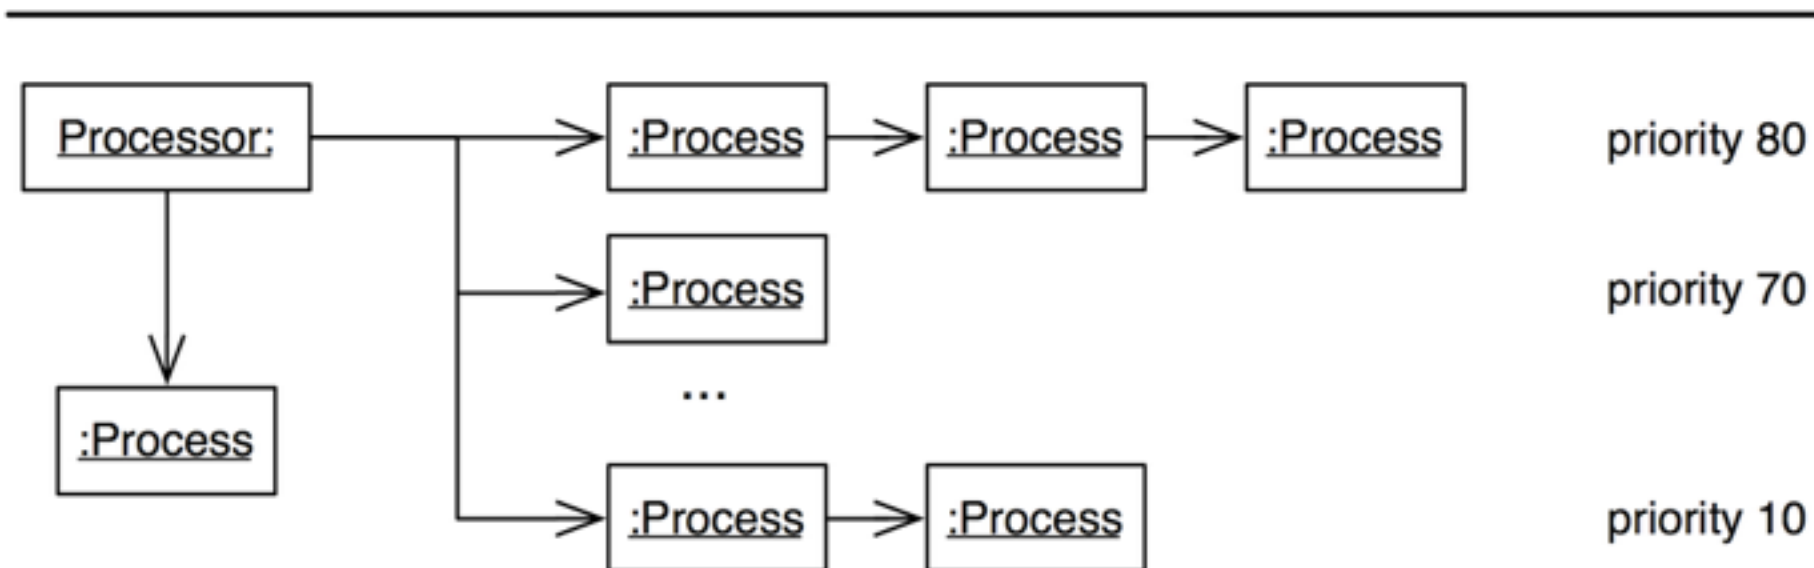- Potentially higher performance
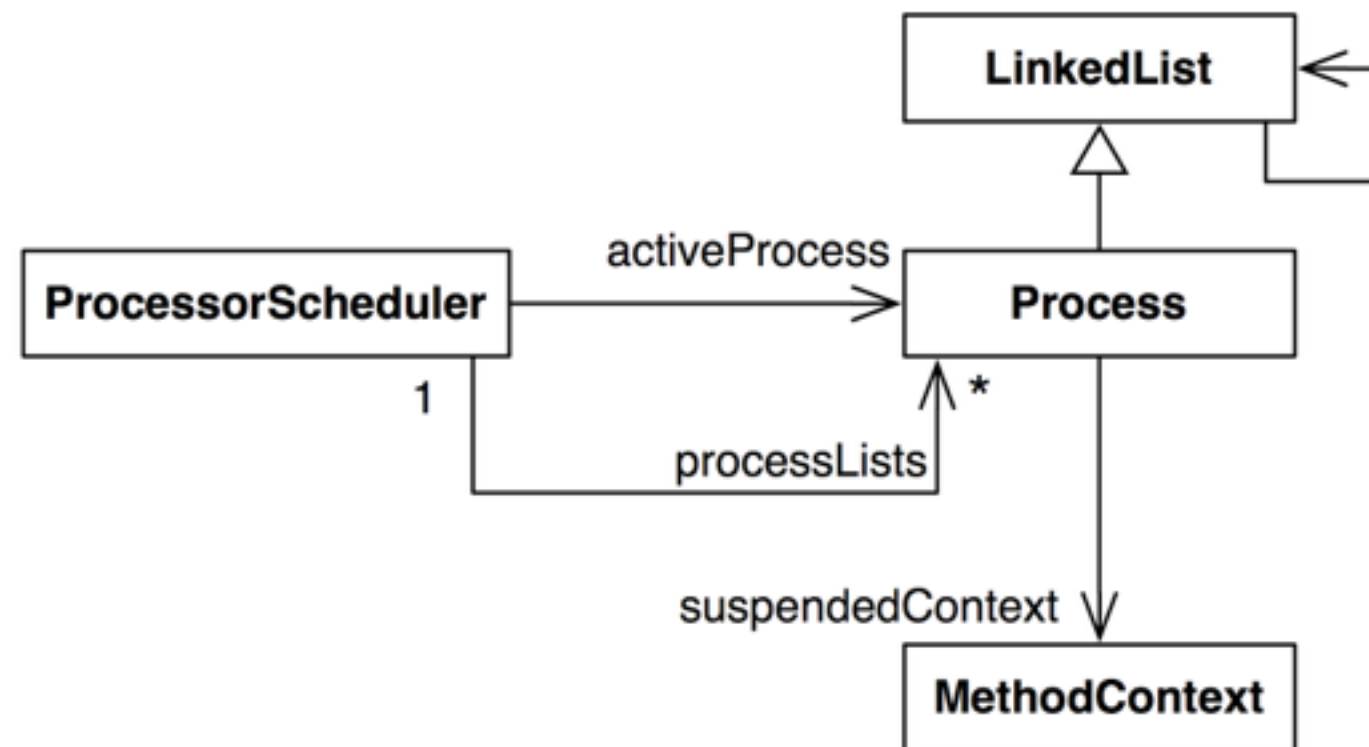
Pharo: Green Threads

Each process has its own execution stack, ip, sp, ...

There is always one (and only one) running process

Each process behaves as if it owns the entire VM

Each process can be interrupted (➜context switching)

Representing Processes and Run Queues

Context Switching

`Interpreter>>transferTo: newProcess`

1. store the current ip and sp registers to the current context
2. store the current context in the old process' `suspendedContext`
3. change `Processor` to point to `newProcess`
4. load ip and sp registers from new process' `suspendedContext`

When you perform a context switch, which process should run next?

Process Scheduler

> *Cooperative* between processes of the same priority

> *Preemptive* between processes of different priorities

Context is switched to the first process with highest priority when:
- current process waits on a semaphore
- current process is suspended or terminated
- `Processor yield` is sent

Context is switched if the following process has a higher priority:
- process is resumed or created by another process
- process is resumed from a signaled semaphore

When a process is interrupted, it moves to the back of its run queue

Example: Semaphores and Scheduling

```
here := false.
lock := Semaphore forMutualExclusion.
[lock critical: [here := true]] fork.
lock critical: [
  self assert: here not.
  Processor yield.
  self assert: here not].
Processor yield.
self assert: here
```

When is the forked process activated?

Roadmap

> Introduction
> Bytecode
> The heap
> Interpreter
> Automatic memory management
> Threading System
> **Optimizations**

Many Optimizations...

> *Method cache* for faster lookup: receiver's class + method selector
> *Method context cache* (as much as 80% of objects created are context objects!)

> *Interpreter loop*: 256 way case statement to dispatch bytecodes

> *Quick returns*: methods that simply return a variable or known constant are compiled as a primitive method

> Small integers are *tagged pointers*: value is directly encoded in field references. Pointer is tagged with low-order bit equal to 1. The remaining 31 bit encode the signed integer value.
> ...

Optimization: JIT (not in Pharo)

Idea: *Just In Time Compilation*

> Translate unit (method, loop, ...) into *native machine code* at runtime

> Store native code in a buffer on the heap

Challenges

> Run-time overhead of compilation

> Machine code takes a lot of space (4-8x compared to bytecode)

> Deoptimization (for debugging) is very tricky

Adaptive compilation: gather statistics to compile only units that are heavily used (*hot spots*)

References

> *Virtual Machines*, Iain D. Craig, Springer, 2006

> *Back to the Future – The Story of Squeak, A Practical Smalltalk Written in Itself*, Ingalls et al, OOPSLA '97

> *Smalltalk-80, the Language and Its Implementation* (AKA "the Blue Book"), Goldberg, Robson, Addison-Wesley, '83
  — http://stephane.ducasse.free.fr/FreeBooks/BlueBook/Bluebook.pdf

> *The Java Virtual Machine Specification*, Second Edition
  — http://java.sun.com/docs/books/jvms/

> *Stacking them up: a Comparison of Virtual Machines*, Gough, IEEE'01

> *Virtual Machine Showdown: Stack Versus Registers*, Shi, Gregg, Beatty, Ertl, VEE'05

What you should know!

✎What is the difference between the operand stack and the execution stack?

✎How do bytecode routines and primitives differ?

✎Why is the object format encoded in a complicated 4bit pattern instead of using regular boolean values?

✎Why is the object address not suitable as a hash value?

✎What happens if an object is only weakly referenced?

✎Why is it hard to build a concurrent mark sweep GC?

✎What does *cooperative multithreading* mean?

✎How do you protect code from concurrent execution?

Can you answer these questions?

✎ There is a lot of similarity between VM and OS design. What are the common components?

✎ Why is accessing the 16th instance variable of an object more efficient than the 17th?

✎ Which disastrous situation could occur if a local C pointer variable exists when a new object is allocated?

✎ Why does #allObjectsDo: not include small integers?

✎ What is the largest possible small integer?

creative commons

Attribution-ShareAlike 4.0 International (CC BY-SA 4.0)

You are free to:

     Share — copy and redistribute the material in any medium or format

     Adapt — remix, transform, and build upon the material for any purpose, even commercially.

     The licensor cannot revoke these freedoms as long as you follow the license terms.

Under the following terms:

Attribution — You must give appropriate credit, provide a link to the license, and indicate if changes were made. You may do so in any reasonable manner, but not in any way that suggests the licensor endorses you or your use.

ShareAlike — If you remix, transform, or build upon the material, you must distribute your contributions under the same license as the original.

No additional restrictions — You may not apply legal terms or technological measures that legally restrict others from doing anything the license permits.

http://creativecommons.org/licenses/by-sa/4.0/